

The Advisor in the Room: AI and the Quiet Loss of Confidants

Monica Hegde
Indiana University Indianapolis
United States
monhegde@iu.edu

Aqueasha Martin-Hammond
Indiana University Indianapolis
United States
aqumarti@iu.edu

Abstract

Most research on AI and human relationships studies companion apps which are designed to simulate a partner or close friend, which can end up taking the place of human partners as people grow attached to them. General-purpose tools like ChatGPT, by contrast, are mostly used in an advisor role, and people already turn to them constantly for something that matters just as much as companionship: as someone to talk to about hard, personal things, like how to handle a fight with a friend or partner. We argue this deserves separate attention, because being the person someone confides in is itself part of what keeps a relationship close, regardless of the advice given. If people bring their hardest moments to AI instead, the people in their lives may quietly lose that role, even without attachment to the AI or less time together. We treat this as a hypothesis, weigh it against two alternative explanations, and outline what it would take to test.

Keywords

Generative AI, Human-AI Interaction, Interpersonal Relationships, Self-Disclosure, AI-Mediated Communication

1 Introduction

Generative AI is increasingly present in moments once reserved for a trusted friend, partner, or family member: working out what to say after a fight, deciding whether a boundary is reasonable, rehearsing a hard conversation before having it. Recent analysis of real conversational AI usage confirms this is not a fringe pattern — people already engage these tools for information seeking, decision support, and emotional support, and interactions are becoming steadily more socially framed over time, with users increasingly treating the AI as an advisor or companion rather than a plain task-completion tool [9]. These are not companionship use cases in the way that term is usually meant — the person is not seeking a relationship with the tool itself — yet they involve exactly the kind of vulnerable self-disclosure that, in human relationships, is understood to do relational work well beyond the immediate exchange of information [1]. The AI acts as a confidant. This paper asks what happens to friendships and marriages when the confidant role, not just the companion role, is increasingly filled by AI. We try to be explicit throughout about which parts of this question existing research already answers, and which parts remain genuinely open and may require more exploration to fully understand the role of AI as a confidant. Distinguishing the confidant role from the companion role has implications for research on GenAI's impact on relationships and communities: it points toward studying relationship-level consequences of AI use rather than attachment or usage patterns alone.

2 The Existing Picture: Companion AI and Attachment Displacement

Engagement with a communication technology can come at the cost of engagement with the people already in a person's life, a pattern originally identified in research on adolescents' online communication and termed the displacement hypothesis: increased use of a technology is associated with a decline in the quality, frequency, or depth of existing offline relationships, whether or not the person perceives this decline as it occurs [15]. Research applying this pattern to AI has focused almost entirely on companion AI — apps explicitly designed to simulate an ongoing romantic or other close social relationships. The core idea in the companion AI literature, sometimes called the social substitution hypothesis, is straightforward: the more time and emotional energy a person invests in an AI companion, the more that AI companion can start to occupy the place a human partner or friend would otherwise hold.

Attachment theory gives this a concrete mechanism. People are thought to organize the important relationships in their lives into something like a ranked hierarchy of attachment figures — who they turn to first, second, or third, when they need support [4, 7]. The claim is that an AI companion that a person becomes emotionally attached to can rise in that hierarchy the same way a human relationship would, and in doing so, can push human partners or friends further down the list, with one documented case of a companion AI user shifting from a counseling service to the chatbot, and another using it as a romantic proxy after an existing relationship had ended [16]. Recent work on adolescents breaks this displacement effect into three more specific forms: a behavioral form (a person actually spends less effort maintaining human relationships), a motivational form (a person still could invest that effort, but no longer feels much need to), and a normative form (a person starts to treat one-sided interaction with an AI as if it were an adequate substitute for a real, reciprocal relationship) [14].

This is a genuinely useful way of thinking about the problem — but it is built entirely around products that are designed to be relational, where the explanatory engine is emotional attachment. The AI has to first become something the person is bonded to before it can push a human relationship out of the way [16]. However, it is not obvious that this framework was built to explain, or even applies to, tools that were never designed to be relational in the first place.

3 What We Know About General-Purpose AI as Confidant

General-purpose GenAI use in this advisory capacity is already documented. An analysis of real ChatGPT conversations found that participants treated the system as an advisor in 47.6% of turns, compared to 38.1% as a plain assistant and only 14.2% as a companion

[9] — showing that, even without any companion-app framing, general-purpose tools are already used heavily for the kind of advisory interaction we describe. Separately, qualitative research on self-care and GenAI has found that people seek advice not only for their own problems but on behalf of friends going through something difficult[5], and other work has specifically examined AI’s role in emotionally loaded moments, such as helping someone draft a breakup message[6].

Two limits on this evidence matter for the argument we want to make. First, social support research distinguishes two kinds of help people seek: emotional support, which involves disclosure sought for comfort or validation in a relational matter, and informational support, which involves practical advice that carries no relational weight. This distinction is well established in chatbot research specifically[10]. The 47.6% advisor-role figure from Karnam et al. (2026) almost certainly blends both kinds together — a question like “how do I talk to my partner about this” alongside one like “how do I file this form.” No existing dataset separates out the emotional-support cases specifically. This means we are extrapolating from a broader category, rather than pointing to data that directly measures the narrower one we care about. Second, and more importantly: this existing research [5, 6, 9] is descriptive. It shows that people already turn to general AI this way, and in some cases it looks closely at a single interaction. But none of it asks the following question: if someone brings a personal topic to AI, are they then less likely to also bring that same topic to a person in their life? That longer-term, relationship-quality question is the actual open gap here — not whether the use pattern exists, which is already well established.

4 Alternative Explanations for the Same Pattern

Social compensation: One potential competing account is that people with smaller or weaker human support networks turn to AI because those relationships were already thin[13], making AI-as-confidant use a symptom of existing relational struggles rather than a cause of their decline. A large survey and chat-log study of AI companion users supports this directly: people with smaller social networks were more likely to use chatbots primarily for companionship, especially when they lacked strong human support [17]. The same study found this companionship-oriented use was linked to lower overall social well-being, particularly with frequent, high-disclosure interaction — a pattern consistent with compensation and displacement operating together rather than one ruling out the other. Our hypothesis does not require compensation to be wrong; both mechanisms could hold for different people, or for the same person at different times, and this alternative already has empirical support in the field[13, 17].

Disclosure rehearsal: A second possibility is that AI use precedes and improves human disclosure rather than replacing it: practicing a hard conversation with AI, drafting a considered message, or getting a calmer outside perspective could make the eventual human conversation happen sooner, go better, or happen at all — consistent with the long-established finding that rehearsing a difficult interaction beforehand improves performance in the real one [12]. A recent study on AI-facilitated preparation for difficult conversations with partners, family, and friends found no evidence

that chatbot rehearsal backfired in participants’ subsequent real conversations, despite concern that AI’s typically agreeable tone might set people up for a mismatch [11]. If this generalizes, AI use would sometimes strengthen human disclosure rather than substitute for it. We cannot rule this out, and future work should test rehearsal against displacement directly — for instance, by tracking whether AI-assisted preparation for a hard topic is followed by the human conversation actually happening, or whether the AI conversation is where the matter quietly ends.

We raise compensation and rehearsal not because they disprove our hypothesis, but because they set a clear bar for it. Displacement requires two things to be true: that confiding in AI usually replaces telling a human rather than preceding or supporting it, and that this pattern is common across many people, not limited mainly to those whose relationships were already weak. If either condition fails, compensation or rehearsal is the better explanation.

5 The Proposed Mechanism: Being Confided In Is Itself Relationship Work

Research on how friendships and relationships grow close treats mutual vulnerability not just as an exchange of information but as one of the actual building blocks of closeness: relationships deepen through a gradual process of increasingly personal, reciprocal disclosure between the two people involved [1]. Being the person someone chooses to bring hard, personal problems to plausibly builds and signals trust on its own, independent of whether the advice that person eventually receives is any good. Whether being displaced as a confidant, specifically, weakens a relationship over time remains an open empirical question, and one we believe deserves direct investigation — which is why we present it here as this paper’s primary central hypothesis.

Consider a concrete case. Suppose a man feels increasingly sidelined by a close friend who has been making plans with their wider friend group without including him, and he is not sure whether to say something or whether he is overreacting. A decade ago, he might have called another close friend (a confidant) to talk it through — venting first, then slowly working toward what he actually wanted to say, with that friend asking questions, pushing back a little, and drawing on years of knowing everyone involved. Today, he might instead open a general-purpose AI tool, describe the situation, and receive calm, well-organized advice: how to raise the concern without sounding accusatory, what to consider before assuming the worst, how to bring it up if he decides to. The advice might be just as good, arguably more measured. But his friend never hears about any of it. Multiply this single substitution across many topics and many months, and what changes has little to do with the quality of advice given: the friend is consulted less often, not because the friendship weakened first, but because the occasion for consulting him no longer arises in the same way.

This is the core distinction our hypothesis rests on. If the mechanism holds, it would not require less time spent together, less affection, or any competition for emotional attachment at all — the two friends could still talk just as often about everything else. It could work through something much quieter: a person simply stops being invited into the vulnerable moments that, added up over time,

may be a large part of what keeps trust between them intact. Neither person would necessarily notice this happening, and neither would experience it as a deliberate choice — the friend is not consciously judged as less important; the occasions to confide in them simply stop arising. Under most existing measures of relationship quality — interaction frequency and closeness [3], or relationship satisfaction [8] — nothing would necessarily look different, because none of those measures ask specifically who someone brought their hardest moments to.

This is also where our hypothesis parts ways with the companion-AI literature, on what, exactly, is being lost. With companion AI, the concern is that a person gets something from the AI itself — comfort, attachment, a feeling of being understood — that competes directly with what a human relationship may provide [14, 16]. Under our hypothesis, the person may walk away with equally good or even better advice from the AI, and may feel no less close to their friend, partner, or sibling in the moment. What the human relationship loses is something no AI conversation can supply no matter how good the advice is: the simple, cumulative fact of having been the one confided in.

The stakes of this hypothesis are not evenly distributed across age groups. Social network research using the convoy model finds that people’s close social ties tend to narrow steadily from adulthood into later life, as retirement, relocation, and the loss of a spouse or long-time friends shrink the pool of relationships available for support [2]. If confiding in AI displaces even one relationship’s role as a confidant, the proportional effect on an older adult, whose remaining convoy of close ties may already be small, is plausibly larger than the same displacement would be for someone with a wider network of peers to draw on. This does not change the mechanism itself, but it does suggest that testing this hypothesis in populations more vulnerable to social declines such as older adults specifically, rather than only in general samples, would reveal where the risk is most concentrated.

6 Implications for Research and Design

Research: Testing this hypothesis against the compensation and rehearsal alternatives calls for longitudinal designs that follow confidant-selection patterns over time, rather than the attachment-strength or usage-time measures common in the companion-AI literature. The central question is not simply whether people use AI as a confidant, this is already established, but whether doing so changes the important human relationships around them. As AI increasingly occupies the confidant role for a given person, does the closeness or quality of their existing friendships and partnerships shift over time, relative to people whose human confiding patterns stay the same? Tracking this longitudinally, rather than through single-topic disclosure snapshots, would allow researchers to distinguish a genuine displacement effect from either compensation [13] (where already-weak relationships preceded AI use) or rehearsal [12] (where AI use left human relationships unchanged or improved).

Design: If this mechanism is real, it raises a design question current GenAI tools have little reason to consider on their own: since the risk lies in AI absorbing the confidant role rather than in the advice itself being poor, could a system recognize categories of disclosure where being confided in by a specific person matters

more than the quality of the advice given? If so, what would it look like to gently point someone back toward a human conversation, and thereby human connections, without presuming to know that person’s relationships better than they do themselves? This is a genuinely difficult design problem — any such feature risks being paternalistic, misjudging what a given relationship actually needs, or simply going unused — but it is worth naming as an open question rather than leaving the default behavior, which is to resolve everything within the conversation itself, unexamined.

7 Conclusion

The companion-AI literature has largely examined whether people form emotional attachment to AI tools. This piece raises a related but distinct question: independent of any attachment to the AI itself, are the people in someone’s life gradually losing their role as confidants, and does that alone carry consequences for relationship quality? We have argued that existing research explores underlying use pattern but has not yet explored this specific mechanism, and we have weighed it directly against social compensation and disclosure rehearsal as competing explanations already present in the literature. We believe this question warrants empirical attention, ahead of the longitudinal evidence that would be needed to settle it.

References

- [1] Irwin Altman and Dalmas A. Taylor. 1973. *Social Penetration: The Development of Interpersonal Relationships*. Holt, Rinehart and Winston, New York.
- [2] Toni C. Antonucci, Kristine J. Ajrouch, and Kira S. Birditt. 2014. The Convoy Model: Explaining Social Relations From a Multidisciplinary Perspective. *The Gerontologist* 54, 1 (2014), 82–92. doi:10.1093/geront/gnt118
- [3] Ellen Berscheid, Mark Snyder, and Allen M. Omoto. 1989. The Relationship Closeness Inventory: Assessing the Closeness of Interpersonal Relationships. *Journal of Personality and Social Psychology* 57, 5 (1989), 792–807.
- [4] John Bowlby. 1969. *Attachment and Loss, Vol. 1: Attachment*. Basic Books, New York.
- [5] Tara Capel, Bernd Ploderer, Filip Bircanin, Simon Hanmer, Jamie Yates, Jiaxuan Wang, Kai Ling Khor, Tuck Wah Leong, Greg Wadley, and Michelle Newcomb. 2024. Studying Self-Care with Generative AI Tools: Lessons for Design. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference*. doi:10.1145/3643834.3661614
- [6] Yue Fu, Yixin Chen, Zelia Gomes Da Costa Lai, and Alexis Hiniker. 2024. Should ChatGPT Write Your Breakup Text? Exploring the Role of AI in Relationship Dissolution. *Proceedings of the ACM on Human-Computer Interaction* 9, CSCW, Article 243 (2024).
- [7] Cindy Hazan and Phillip Shaver. 1987. Romantic Love Conceptualized as an Attachment Process. *Journal of Personality and Social Psychology* 52, 3 (1987), 511–524. doi:10.1037/0022-3514.52.3.511
- [8] Susan S. Hendrick. 1988. A Generic Measure of Relationship Satisfaction. *Journal of Marriage and the Family* 50, 1 (1988), 93–98.
- [9] Sai Keerthana Karnam et al. 2026. Bowling with ChatGPT: On the Evolving User Interactions with Conversational AI Systems. In *Proceedings of the ACM Web Conference 2026*. doi:10.1145/3774904.3792978
- [10] Inju Lee and Sowon Hahn. 2024. On the Relationship Between Mind Perception and Social Support of Chatbots. *Frontiers in Psychology* 15, Article 1282036 (2024). doi:10.3389/fpsyg.2024.1282036
- [11] Charul Maheshka, Dunigan Folk, and Elizabeth Dunn. 2026. Harnessing Artificial Intelligence to Facilitate Difficult Conversations. *Social Psychological and Personality Science* (2026). doi:10.1177/19485506261444068
- [12] Richard M. McFall and Albert R. Marston. 1970. An Experimental Investigation of Behavior Rehearsal in Assertive Training. *Journal of Abnormal Psychology* 76, 2 (1970), 295–303.
- [13] Katelyn Y. A. McKenna and John A. Bargh. 2000. Plan 9 from Cyberspace: The Implications of the Internet for Personality and Social Psychology. *Personality and Social Psychology Review* 4, 1 (2000), 57–75.
- [14] Xinyue Sun, Yifan Wang, and Brandon T. McDaniel. 2026. AI Companions and Adolescent Social Relationships: Benefits, Risks, and Bidirectional Influences. *Child Development Perspectives* 20, 2 (2026), 91–98.

- [15] Patti M. Valkenburg and Jochen Peter. 2007. Online Communication and Adolescent Well-Being: Testing the Stimulation Versus the Displacement Hypothesis. *Journal of Computer-Mediated Communication* 12, 4 (2007), 1169–1182. doi:10.1111/j.1083-6101.2007.00368.x
- [16] Tianling Xie and Iryna Pentina. 2022. Attachment Theory as a Framework to Understand Relationships with Social Chatbots: A Case Study of Replika. In *Proceedings of the 55th Hawaii International Conference on System Sciences*. 2046–2055. doi:10.24251/HICSS.2022.258
- [17] Yutong Zhang, Dora Zhao, Jeffrey T. Hancock, Robert Kraut, and Diyi Yang. 2025. The Rise of AI Companions: How Human-Chatbot Relationships Influence Well-Being. *arXiv preprint arXiv:2506.12605* (2025). <https://arxiv.org/abs/2506.12605>